



Introduction to hybrid modeling

Basil Kraft

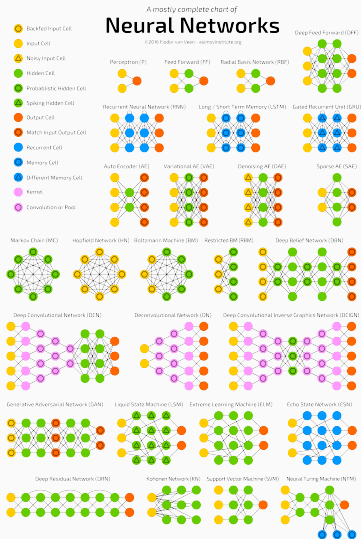
February 2023

MPI for Biogeochemistry
Department Biogeochemical Integration
Jena

- Quick recap:
 - Neural networks
 - Demo: neural networks with PyTorch
- Modeling snow with a hybrid model
 - Background & data simulation
 - Demo: hybrid modeling


Artificial neural networks (ANNs)

Artificial neural networks (ANNs)



- ANNs are **highly flexible**; We can adapt the model architecture to our needs / to fit the problem
- ANN's performance **scales with data**
- Computation based on (many) matrix operations; **computation is scaleable**

Artificial neural networks (ANNs)



- ANNs consist of (learned) **sequentially arranged nonlinear transformations**
- Each transformation is simple, but in combination, ANNs are **highly expressive**

Figure 1: ANNs are hierarchical feature learners.

Artificial neural networks (ANNs)





Figure 2: The perceptron: simple architecture for binary classification (or regression with target domain $-1 < y < 1$)

- The building blocks of ANNs are simple
- So-called **activation functions**^a add non-linearity

^aa non-linear, differentiable function

Artificial neural networks (ANNs)



- The building blocks of ANNs are simple
- So-called **activation functions**^a add non-linearity

^aa non-linear, differentiable function

Figure 3: ANNs are hierarchical feature learners.




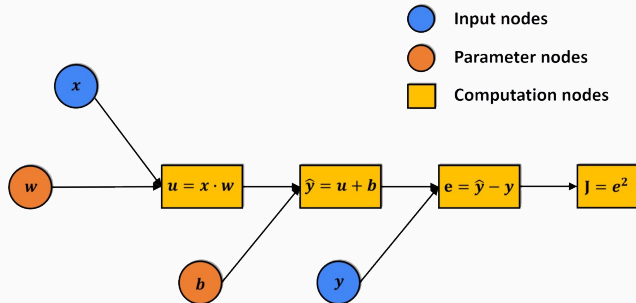
Figure 4: A 2-D loss landscape: the z-axis is the loss (which we seek to minimize), x and y are parameter space.

- ANNs are optimized with gradient descent
- The parameters are adapted iteratively by changing them in little “downhill” steps

Artificial neural networks (ANNs)

- Example: a simple linear model with one input and one output
- How do we efficiently optimize this?


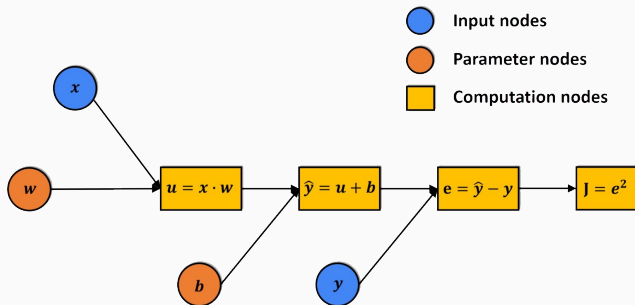
$$\frac{\partial J}{\partial w} = ?$$
$$\frac{\partial J}{\partial b} = ?$$



Artificial neural networks (ANNs)

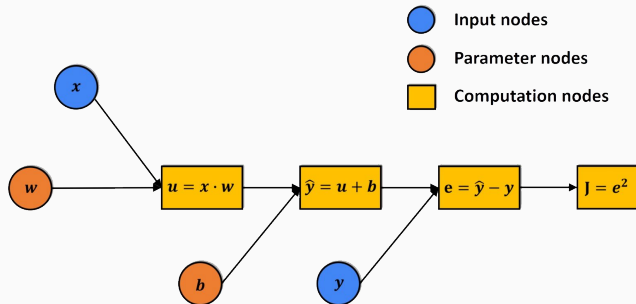
- Example: a simple linear model with one input and one output
- How do we efficiently optimize this?

$$\frac{\partial J}{\partial w} = ?$$
$$\frac{\partial J}{\partial b} = ?$$




Artificial neural networks (ANNs)


- We apply the chain rule: $\frac{\partial a}{\partial c} = \frac{\partial a}{\partial b} \frac{\partial b}{\partial c}$
- Blue terms have to be calculated only once!



$$\frac{\partial J}{\partial w} = \frac{\partial J}{\partial e} \frac{\partial e}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial w}$$
$$\frac{\partial J}{\partial b} = \frac{\partial J}{\partial e} \frac{\partial e}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b}$$



Artificial neural networks (ANNs)



- The error is **backpropagated** layer-wise and summed up at nodes
- Backpropagation is agnostic of the full model structure, we only need to know the local gradient computation.


Artificial neural networks (ANNs)

Pseudo-code for neural net optimization


- Multiple iterations through training data (epochs)
- In each epoch, we iterate the data in minibatches without replacement

```
1: initialize(model.params)
2: for  $e$  in  $Epochs$  do
3:   for  $batch$  in  $Data$  do
4:      $zero\_grads(model.params.grads)$  # reset gradients
5:      $x, y = getXY(batch)$  # get batch
6:      $\hat{y} = model(x)$  # model forward pass
7:      $loss = f_{loss}(\hat{y}, y)$  # loss calculation
8:      $model.params.grads = backprop(loss)$  # compute gradients
9:      $update(model.params)$ 
10:  end for
11: end for
```


Gradient Descent



Stochastic Gradient Descent



Mini-Batch Gradient Descent



- ANNs are extremely powerful (universal function approximators)

Artificial neural networks (ANNs)

- ANNs are extremely powerful (universal function approximators)
- Using large models requires lots of training data

- ANNs are extremely powerful (universal function approximators)
- Using large models requires lots of training data
- We can make models more data-efficient by using tailored architectures

Artificial neural networks (ANNs)

- ANNs are extremely powerful (universal function approximators)
- Using large models requires lots of training data
- We can make models more data-efficient by using tailored architectures
- Careful with overfitting

- ANNs are extremely powerful (universal function approximators)
- Using large models requires lots of training data
- We can make models more data-efficient by using tailored architectures
- Careful with overfitting
- Luckily, there are frameworks for building and training ANNs (PyTorch, Tensorflow, etc.)


Demo: neural networks with PyTorch

Hybrid modeling

Hybrid model of snow water equivalent

We model snow water equivalent S as a function of precipitation p , air temperature T , and solar irradiation r :

$$S_t = f(p_t, T_t, r_t)$$




Hybrid model of snow water equivalent

We use a neural network to learn the process from data:

$$S_t = f_{\text{NN}}(p_t, T_t, r_t)$$


This may work quite well, but

- The model can predict implausible values (e.g., $S < 0$, or snowfall but no precipitation)
- We have no control over the process, no prior knowledge used
- We cannot physically interpret what the model learned
- S is a state we can't predict from just today's weather
- **We can do better**



Restricting the output space such that $S \geq 0$

$$S_t = \text{Sortplus}(f_{\text{NN}}(p_t, T_t, r_t))$$



Account for memory effects


$$S_t = \text{Sortplus}(f_{\text{NN}}(p_{k \leq t}, T_{k \leq t}, r_{k \leq t}))$$



→ standard practices, not hybrid modeling

Implement explicit process knowledge (hybrid modeling)

- We assume that two main processes contribute to ΔS (we simplify a bit)
 - snowmelt s_{melt}
 - snowfall s_{fall}
- Update snow with the two quantities:
 - $$S_t = S_{t-1} + \overbrace{\text{Sortplus}(f_{\text{NN}}^1(p_t, T_t, r_t))}^{s_{\text{fall},t}} - \overbrace{\text{Sortplus}(f_{\text{NN}}^2(p_t, T_t, r_t))}^{s_{\text{melt},t}}$$
- Is this a good idea? Why (not)?



- Snow fall without precipitation possible
- Snow melt with temperatures below freezing point
- Negative snow possible
- Equifinality: $S_t = S_{t-1} + S_{\text{fall},t} - S_{\text{melt},t}$
→ model not very useful; we need better constraints!





Figure 5: Equifinality: different solutions lead to the same result


Hybrid modeling

$$\bullet s_{\text{fall},t} = p_t \underbrace{\alpha_{\text{fall},t}}^{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$$




Hybrid modeling

- $$s_{\text{fall},t} = p_t \underbrace{\alpha_{\text{fall},t}}^{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$$
- $$s_{\text{melt},t} = \max(0, T_t) \underbrace{\alpha_{\text{melt},t}}^{f_{\text{NN}}^2(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{melt},t}$$




Hybrid modeling

- $$S_{\text{fall},t} = p_t \underbrace{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}_{\alpha_{\text{fall},t}}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$$
- $$S_{\text{melt},t} = \max(0, T_t) \underbrace{f_{\text{NN}}^2(p_t, T_t, r_t, S_{t-1})}_{\alpha_{\text{melt},t}}, \quad 0 \leq \alpha_{\text{melt},t}$$
- $$S_t = \max(0, S_{t-1} + S_{\text{fall},t} - S_{\text{melt},t}) \quad (S_{\text{melt},t} \text{ is actually potential snow melt here})$$




Hybrid modeling

- $$S_{\text{fall},t} = p_t \underbrace{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}_{\alpha_{\text{fall},t}}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$$
- $$S_{\text{melt},t} = \max(0, T_t) \underbrace{f_{\text{NN}}^2(p_t, T_t, r_t, S_{t-1})}_{\alpha_{\text{melt},t}}, \quad 0 \leq \alpha_{\text{melt},t}$$
- $$S_t = \max(0, S_{t-1} + S_{\text{fall},t} - S_{\text{melt},t}) \quad (S_{\text{melt},t} \text{ is actually } \text{potential snow melt here})$$
- Only snow fall with $p > 0$, only snow melt with $T > 0$




Hybrid modeling

- $S_{\text{fall},t} = p_t \underbrace{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}_{\alpha_{\text{fall},t}}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$
- $S_{\text{melt},t} = \max(0, T_t) \underbrace{f_{\text{NN}}^2(p_t, T_t, r_t, S_{t-1})}_{\alpha_{\text{melt},t}}, \quad 0 \leq \alpha_{\text{melt},t}$
- $S_t = \max(0, S_{t-1} + S_{\text{fall},t} - S_{\text{melt},t})$ ($S_{\text{melt},t}$ is actually **potential** snow melt here)
- Only snow fall with $p > 0$, only snow melt with $T > 0$
- No negative snow




Hybrid modeling

- $S_{\text{fall},t} = p_t \underbrace{\alpha_{\text{fall},t}}_{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$
- $S_{\text{melt},t} = \max(0, T_t) \underbrace{\alpha_{\text{melt},t}}_{f_{\text{NN}}^2(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{melt},t}$
- $S_t = \max(0, S_{t-1} + S_{\text{fall},t} - S_{\text{melt},t})$ ($S_{\text{melt},t}$ is actually **potential** snow melt here)
- Only snow fall with $p > 0$, only snow melt with $T > 0$
- No negative snow
- The neural nets are informed about the current snow state as we pass S_{t-1} as input.




Hybrid modeling

- $S_{\text{fall},t} = p_t \underbrace{\alpha_{\text{fall},t}}_{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$
- $S_{\text{melt},t} = \max(0, T_t) \underbrace{\alpha_{\text{melt},t}}_{f_{\text{NN}}^2(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{melt},t}$
- $S_t = \max(0, S_{t-1} + S_{\text{fall},t} - S_{\text{melt},t})$ ($S_{\text{melt},t}$ is actually **potential** snow melt here)
- Only snow fall with $p > 0$, only snow melt with $T > 0$
- No negative snow
- The neural nets are informed about the current snow state as we pass S_{t-1} as input.
- We gain interpretability of snow melt and snow fall, which we do not have with plain neural network approach




Hybrid modeling

- $S_{\text{fall},t} = p_t \underbrace{\alpha_{\text{fall},t}}_{f_{\text{NN}}^1(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{fall},t} \leq 1$
- $S_{\text{melt},t} = \max(0, T_t) \underbrace{\alpha_{\text{melt},t}}_{f_{\text{NN}}^2(p_t, T_t, r_t, S_{t-1})}, \quad 0 \leq \alpha_{\text{melt},t}$
- $S_t = \max(0, S_{t-1} + S_{\text{fall},t} - S_{\text{melt},t})$ ($S_{\text{melt},t}$ is actually **potential** snow melt here)
- Only snow fall with $p > 0$, only snow melt with $T > 0$
- No negative snow
- The neural nets are informed about the current snow state as we pass S_{t-1} as input.
- We gain interpretability of snow melt and snow fall, which we do not have with plain neural network approach
- What are potential pitfalls?




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:
 - *E.g.*, snow fall includes wind transport from neighboring areas




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:
 - *E.g.*, snow fall includes wind transport from neighboring areas
 - *E.g.*, snow melt includes sublimation




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:
 - E.g., snow fall includes wind transport from neighboring areas
 - E.g., snow melt includes sublimation
- Biases due to **wrong assumptions**:
 - E.g., we could have snow melt with $T < 0$ because T is air temperature.




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:
 - E.g., snow fall includes wind transport from neighboring areas
 - E.g., snow melt includes sublimation
- Biases due to **wrong assumptions**:
 - E.g., we could have snow melt with $T < 0$ because T is air temperature.
- Biases due to **data**:




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:
 - E.g., snow fall includes wind transport from neighboring areas
 - E.g., snow melt includes sublimation
- Biases due to **wrong assumptions**:
 - E.g., we could have snow melt with $T < 0$ because T is air temperature.
- Biases due to **data**:
 - What if precipitation is underestimated? (actually the case for snowfall)




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:
 - E.g., snow fall includes wind transport from neighboring areas
 - E.g., snow melt includes sublimation
- Biases due to **wrong assumptions**:
 - E.g., we could have snow melt with $T < 0$ because T is air temperature.
- Biases due to **data**:
 - What if precipitation is underestimated? (actually the case for snowfall)
 - What if air temperature is overestimated?




Hybrid modeling

- Equifinality if $p > 0$ & $T > 0$?
- Biases due to **neglected processes**:
 - E.g., snow fall includes wind transport from neighboring areas
 - E.g., snow melt includes sublimation
- Biases due to **wrong assumptions**:
 - E.g., we could have snow melt with $T < 0$ because T is air temperature.
- Biases due to **data**:
 - What if precipitation is underestimated? (actually the case for snowfall)
 - What if air temperature is overestimated?
 - What if snow is biased? (many observation-based snow products saturate at higher levels)



How do we implement hybrid models?



- There is no qualitative difference between neural network layers and physical equations
- The physical equations need to be differentiable and that's it, backpropagation still works!



Demo: hybrid modeling of snow water equivalent with PyTorch

Artificial neural networks (ANNs)

- Daily data from ERA5 reanalysis
- Training/validation split in time (to keep things simple)



- Hybrid models can be implemented in deep learning frameworks such as PyTorch
- Model design, data, and assumptions play a key role
- Hybrid modeling has some strengths, but also weaknesses
- Compared to machine learning, hybrid models can be physically more plausible and to a certain extent interpretable